

# Learning Discriminative Representations to Interpret Image Recognition Models

Thèse de Doctorat

Felipe Torres Figueroa

École Centrale de Marseille

Laboratoire d'Informatique et de Systèmes (LIS)

Marseille, Septembre 23 2024

*Membres de jury*

Président: Frédéric Jurie

Examineurs: Frédéric Jurie  
Giorgos Tolias

Rapporteurs: Frédéric Precioso  
Diane Larlus

Superviseurs: Stéphane Ayache  
Ronan Sicre  
Yannis Avrithis

# Table of Contents

- Introduction
- 1 Background
- 2 Opti-CAM: Optimizing saliency maps for interpretability
- 3 CA-Stream: Attention-based pooling for interpretable image recognition
- 4 A learning paradigm for interpretable gradients
- References


# Table of Contents

- Introduction
- 1 Background
- 2 Opti-CAM: Optimizing saliency maps for interpretability
- 3 CA-Stream: Attention-based pooling for interpretable image recognition
- 4 A learning paradigm for interpretable gradients
- References

# Motivation

## High stakes

Consider these situations:



# Motivation

High stakes


Consider these situations:



# Motivation

High stakes

Consider these situations:



## Straight to the point



- How do we **know how** a system works?
- How do we **know how** safe a system is?
- If a system fails, **who** is accountable?

We must **understand** the behaviour of these models.

## Step by step

# Computation, Computer Vision and AI



## Explainable AI



## Thesis objectives

# Computation



Better known as *Computer Science*.

## Study of:

- Algorithms.
- Data structures.
- Design of hardware and software.




# Computation, Computer Vision and AI

## Computer Vision

Replication of human vision capabilities.

Three fundamental tasks[1]:

- Recognition.
- Reconstruction.
- Reorganization.



# Computation, Computer Vision and AI

## Artificial Intelligence

Systems capable of performing  
tasks requiring human intelligence  
[2].


# Computation, Computer Vision and AI

## Artificial Intelligence

Systems capable of performing tasks requiring human intelligence [2].

Subfields:

- Machine Learning (*ML*) & Deep Learning (*DL*).
- Computer Vision(*CV*).
- Natural Language Processing (*NLP*).
- Robotics.



# Explainable AI

We are interested in understanding models, behaving like a black box model:



# Explainable AI


We are interested in understanding models, behaving like a black box model:



**We want to *know why*  $f(x) \rightarrow y$**


# Explainable AI

Desired data of Interpretability Study [4]




Regulation planned with the European AI act[3]


# Fitting it all together



# Fitting it all together



# Fitting it all together



# Thesis Objectives

**Improvement of recognition and interpretable properties of model predictions.**

# Thesis Objectives

## **Improvement of recognition and interpretable properties of model predictions.**

In particular:

- Development of low cost/complexity explainability approaches.
- Establishment of a fixed evaluation protocol.
- Differentiation of human based and machine explanations.

# Table of Contents

- Introduction
- 1 Background
- 2 Opti-CAM: Optimizing saliency maps for interpretability
- 3 CA-Stream: Attention-based pooling for interpretable image recognition
- 4 A learning paradigm for interpretable gradients
- References





# Image Recognition Models

## Convolutional Neural Networks (CNNs)

Based on the **convolution operation**.

A representation  $f \star g$  is computed for a feature map  $f$  and a kernel  $g$ .



# Image Recognition Models

## Convolutional Neural Networks (CNNs)

Training deep models is difficult, gradient information *is* lost with depth.

A model surviving the test of time: *ResNet* [5].


# Image Recognition Models

## Convolutional Neural Networks (CNNs)


Training deep models is difficult, gradient information *is* lost with depth.

A model surviving the test of time: *ResNet* [5].


Input image



### Basic Block



### Bottleneck Block



# Image Recognition Models

## Self Attention Architectures

### Transformers

*Attention:* compute the relevance of each element relative to others in an embedding.

$$\text{Self Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V$$

(1)

# Image Recognition Models

## Self Attention Architectures

# Transformers


*Attention:* compute the relevance of each element relative to others in an embedding.

$$\text{Self Attention}(Q, K, V) = \text{softmax} \left( \frac{QK^T}{\sqrt{d_k}} \right) V \quad (1)$$

# Image Recognition Models

## Self Attention Architectures



An adoption of transformers, the *Vision Transformer* (ViT) [7]



# Image Recognition Models

## Self Attention Architectures


An adoption of transformers, the *Vision Transformer* (ViT) [7]



# Image Recognition Models

## The Current Landscape

Transformers had a strong impact on image recognition:



# Interpretability

## Two dimension grouping

### Transparency

Approaches requiring changes to an architecture or training process.

### Post-Hoc Interpretability

Approaches utilizing information contained within a model.

# Interpretability

## Class Activation Methods

Semantic information is found in deeper stages of a model [8].




# Interpretability

## Class Activation Methods

Semantic information is found in deeper stages of a model [8].



Salient regions are highlighted for an input image:



Governed by equation:

$$S_\ell^c(\mathbf{u}) := h \left( \sum_k w_k^c A_\ell^k \right) \quad (2)$$

# Interpretability

## Evaluating Interpretability

### Image Classification Approach

An explanation should demonstrate similar predictive properties to its query:

Input image. ( $I$ )



$$P_i^c = 0.8756$$

Explanation Map. ( $E^c$ )



$$O_i^c = 0.7442$$

# Interpretability

## Evaluating Interpretability

### Image Classification Approach

An explanation should demonstrate similar predictive properties to its query:

Input image. ( $I$ )



$$P_i^c = 0.8756$$

Explanation Map. ( $E^c$ )



$$O_i^c = 0.7442$$

- Prediction loss when using an explanation. (Average Drop (AD) [9])
- Ratio of explanations being better than the input image. (Average Increase (AI) [9])
- Reaction to saliency guided modifications (Insertion/Deletion (Ins/Del) [10])

# Interpretability

## Evaluating Interpretability

### Weakly Supervised Object Localization Approach

A saliency map should be good to localize the object of interest for an example:

Input image. ( $I$ )



Saliency Map. ( $E^c$ )



# Interpretability

## Evaluating Interpretability

### Weakly Supervised Object Localization Approach

A saliency map should be good to localize the object of interest for an example:

Input image. ( $I$ )



Saliency Map. ( $E^c$ )



- Relationships between the bounding boxes (OM, LE,  $F_1$ ).
- Overlap between object mask and bounding box (BA).
- Localization of most salient points and fraction of them in the bounding box (EP, SP).
- Penalty of the size of the bounding box and image size, regarding prediction (SM).


# Table of Contents

- Introduction
- 1 Background
- 2 Opti-CAM: Optimizing saliency maps for interpretability**
- 3 CA-Stream: Attention-based pooling for interpretable image recognition
- 4 A learning paradigm for interpretable gradients
- References

# Preliminaries

## Premise

*Is it possible to optimize the weighting coefficient  $w_k^c$  to obtain a saliency map with the highest probability of prediction?*



# Preliminaries

## Contributions

### Opti-CAM

Maximization of predicted probability of an explanation.

### Average Drop

Evaluation metric designed to enhance comparisons.

# Opti-CAM

## Motivation

- **Combination of Feature Maps**

Feature maps can be considered independent to compute  $w_k^c$  [11]. *What about a linear combination?*

$$F(\mathbf{w}) := f \left( \mathbf{u} \odot n \left( \text{up} \left( \sum_k w_k A_l^k \right) \right) \right)_c$$

# Opti-CAM

## Motivation

- **Combination of Feature Maps**

Feature maps can be considered independent to compute  $w_k^c$  [11]. *What about a linear combination?*

$$F(\mathbf{w}) := f \left( \mathbf{u} \odot n \left( \text{up} \left( \sum_k w_k A_l^k \right) \right) \right)_c$$

- **Masking-based Saliency** Saliency maps produced by optimization in the input space [12, 13]  
The optimization usually takes form of:

$$S^c(\mathbf{u}) := \underset{\mathbf{m} \in M}{\operatorname{argmax}} f(\mathbf{u} \odot n(\text{up}(\mathbf{m})))_c + \delta R(\mathbf{m})$$

# Opti-CAM

## Optimization

**Formulation** Assuming channel weights  $w_k := \text{softmax}(\mathbf{u})_k$ , the saliency map  $S_\ell$  is a function of both an input image  $x$  and a variable  $\mathbf{u}$ :

$$S_\ell(\mathbf{x}; \mathbf{u}) := \sum_k \text{softmax}(\mathbf{u})_k A_\ell^k \quad (3)$$

# Opti-CAM

## Optimization

**Formulation** Assuming channel weights  $w_k := \text{softmax}(\mathbf{u})_k$ , the saliency map  $S_\ell$  is a function of both an input image  $x$  and a variable  $\mathbf{u}$ :

$$S_\ell(\mathbf{x}; \mathbf{u}) := \sum_k \text{softmax}(\mathbf{u})_k A_\ell^k \quad (3)$$

Find vector  $\mathbf{u}^*$  maximizing confidence for class  $c$ , with an input  $\mathbf{x}$  masked by Saliency Map  $S_\ell(\mathbf{x}; \mathbf{u}^*)$ :

$$\mathbf{u}^* := \underset{\mathbf{u}}{\operatorname{argmax}} F_\ell^c(\mathbf{x}; \mathbf{u}^*), \quad (4)$$

# Opti-CAM

## Optimization

**Formulation** Assuming channel weights  $w_k := \text{softmax}(\mathbf{u})_k$ , the saliency map  $S_\ell$  is a function of both an input image  $x$  and a variable  $\mathbf{u}$ :

$$S_\ell(\mathbf{x}; \mathbf{u}) := \sum_k \text{softmax}(\mathbf{u})_k A_\ell^k \quad (3)$$

Find vector  $\mathbf{u}^*$  maximizing confidence for class  $c$ , with an input  $\mathbf{x}$  masked by Saliency Map  $S_\ell(\mathbf{x}; \mathbf{u}^*)$ :

$$\mathbf{u}^* := \underset{\mathbf{u}}{\operatorname{argmax}} F_\ell^c(\mathbf{x}; \mathbf{u}^*), \quad (4)$$

The objective function:

$$F_\ell^c(\mathbf{x}; \mathbf{u}) := g_c(f(x \odot S_\ell(\mathbf{x}; \mathbf{u}))) \quad (5)$$

# Opti-CAM

## Optimization

**Formulation** Assuming channel weights  $w_k := \text{softmax}(\mathbf{u})_k$ , the saliency map  $S_\ell$  is a function of both an input image  $x$  and a variable  $\mathbf{u}$ :

$$S_\ell(\mathbf{x}; \mathbf{u}) := \sum_k \text{softmax}(\mathbf{u})_k A_\ell^k \quad (3)$$

Find vector  $\mathbf{u}^*$  maximizing confidence for class  $c$ , with an input  $\mathbf{x}$  masked by Saliency Map  $S_\ell(\mathbf{x}; \mathbf{u}^*)$ :

$$\mathbf{u}^* := \underset{\mathbf{u}}{\operatorname{argmax}} F_\ell^c(\mathbf{x}; \mathbf{u}), \quad (4)$$

The objective function:

$$F_\ell^c(\mathbf{x}; \mathbf{u}) := g_c(f(\mathbf{x} \odot S_\ell(\mathbf{x}; \mathbf{u}))) \quad (5)$$

Then:

$$S_\ell^c(\mathbf{x}) := S_\ell(\mathbf{x}; \mathbf{u}^*) = S_\ell(\mathbf{x}; \underset{\mathbf{u}}{\operatorname{argmax}} F_\ell^c(\mathbf{x}; \mathbf{u})) \quad (6)$$

# Average Gain

## Motivation

Existing metrics for measuring interpretability ( $AD_{\downarrow}$ ,  $AI_{\uparrow}$ ), can be fooled.

# Average Gain

## Motivation

Existing metrics for measuring interpretability ( $AD\downarrow$ ,  $AI\uparrow$ ), can be fooled.

All pixels are salient *but-one*:

Input



$$p_{389} = 0.874$$

Explanation



$$o_{389} = 0.822$$

- Drop = 5.95%
- Increase = 0%

# Average Gain

## Definition

With the notation on slide (22):

$$\text{AG}(\%) := \frac{1}{N} \sum_{i=1}^N \frac{[o_i^c - p_i^c]_+}{1 - p_i^c} \cdot 100 \quad (7)$$

Measuring the increase in predictive power when masking an input.

# Experimental Set-Up

## Implementation

- All images are set to  $224 \times 224 \times 3$ .
- Adam [14] optimizer, for 100 iterations.
- Cross Entropy.
- Stop optimizing if loss is less than  $10^{-10}$ .

---

<sup>1</sup>Pytorch Model Zoo <https://pytorch.org/vision/0.8/models.html>

<sup>2</sup>Timm <https://github.com/rwightman/pytorch-image-models>

# Experimental Set-Up

## Implementation

- All images are set to  $224 \times 224 \times 3$ .
- Adam [14] optimizer, for 100 iterations.
- Cross Entropy.
- Stop optimizing if loss is less than  $10^{-10}$ .

## Datasets and Models


- **ILSVRC** Validation splits of ImageNet ([15, 16]), 50k images, over 1k classes.
- **Medical Data** Chest X-ray for 3 diseases [17]. Endoscopy data for pathologies, findings and landmarks [18].
- **Networks** ResNet-50 [5], and VGG [19]. <sup>1</sup>ViT-B (16-224) [7] and DeiT-B (16-224) [20] <sup>2</sup>.

---

<sup>1</sup>Pytorch Model Zoo <https://pytorch.org/vision/0.8/models.html>


<sup>2</sup>Timm <https://github.com/rwightman/pytorch-image-models>  

## Qualitative Evaluation




# Results

## Qualitative Evaluation



# Results

## Qualitative Evaluation



# Results

## Quantitative Evaluation

### Classification Experiments:

| METHOD         | RESNET50   |             |             |             |             |        | VGG16      |             |             |             |            |       |
|----------------|------------|-------------|-------------|-------------|-------------|--------|------------|-------------|-------------|-------------|------------|-------|
|                | AD ↓       | AG ↑        | AI ↑        | I ↑         | D ↓         | T(s)   | AD ↓       | AG ↑        | AI ↑        | I ↑         | D ↓        | T(s)  |
| Fake-CAM       | 0.8        | 1.6         | 46.0        | 50.7        | 28.1        | 0.00   | 0.5        | 0.6         | 42.6        | 46.1        | 26.9       | 0.00  |
| Grad-CAM       | 12.2       | 17.6        | 44.4        | 66.3        | 14.7        | 0.03   | 14.2       | 14.7        | 40.6        | <b>64.1</b> | 11.6       | 0.02  |
| Grad-CAM++     | 12.9       | 16.0        | 42.1        | 66.0        | 14.7        | 0.03   | 17.1       | 10.2        | 33.4        | 62.9        | 12.2       | 0.02  |
| Score-CAM      | 8.6        | 26.6        | 56.7        | 65.7        | 16.3        | 15.22  | 13.5       | 15.6        | 41.7        | 62.5        | 12.1       | 3.11  |
| Ablation-CAM   | 12.5       | 16.4        | 42.8        | 65.9        | 14.6        | 18.26  | 15.5       | 12.6        | 36.9        | 63.8        | 11.4       | 2.98  |
| XGrad-CAM      | 12.2       | 17.6        | 44.4        | 66.3        | 14.7        | 0.03   | 13.8       | 14.8        | 41.2        | <b>64.1</b> | 11.7       | 0.02  |
| Layer-CAM      | 15.6       | 15.0        | 38.8        | 67.0        | <b>14.2</b> | 0.08   | 48.9       | 3.1         | 13.5        | 58.3        | <b>6.4</b> | 0.07  |
| ExPerturbation | 38.1       | 9.5         | 22.5        | <b>70.7</b> | 15.0        | 152.96 | 43.0       | 7.1         | 20.5        | 61.1        | 15.0       | 83.20 |
| Opti-CAM       | <b>1.5</b> | <b>68.8</b> | <b>92.8</b> | 62.0        | 19.7        | 4.15   | <b>1.3</b> | <b>71.2</b> | <b>92.7</b> | 59.2        | 11.0       | 3.94  |

# Results

## Quantitative Evaluation

### Localization Experiments:

| METHOD         | RESNET50    |             |             |             |             |             |             |  | VGG16       |             |             |             |             |             |             |  |
|----------------|-------------|-------------|-------------|-------------|-------------|-------------|-------------|--|-------------|-------------|-------------|-------------|-------------|-------------|-------------|--|
|                | OM↓         | LE↓         | F1↑         | BA↑         | SP↑         | EP↑         | SM↓         |  | OM↓         | LE↓         | F1↑         | BA↑         | SP↑         | EP↑         | SM↓         |  |
| Fake-CAM       | 63.6        | 54.0        | 57.7        | 47.9        | 99.8        | 28.5        | 0.98        |  | 64.7        | 54.0        | 57.7        | 47.9        | 99.8        | 28.5        | 1.07        |  |
| Grad-CAM       | 72.9        | 65.8        | 49.8        | <b>56.2</b> | 69.8        | 33.3        | 1.30        |  | 71.1        | 62.3        | 42.0        | 54.2        | 64.8        | 32.0        | 1.39        |  |
| Grad-CAM++     | 73.1        | 66.1        | <b>50.4</b> | <b>56.2</b> | 69.9        | 33.1        | 1.29        |  | 70.8        | 61.9        | 44.3        | 55.2        | 66.2        | 32.3        | 1.38        |  |
| Score-CAM      | <b>72.2</b> | 64.9        | 49.6        | 54.5        | 68.7        | 32.4        | <b>1.25</b> |  | 71.2        | 62.5        | <b>45.3</b> | <b>58.5</b> | <b>68.2</b> | 33.4        | 1.40        |  |
| Ablation-CAM   | 72.8        | 65.7        | 50.2        | 56.1        | 69.9        | 33.1        | 1.26        |  | 71.3        | 62.6        | 43.2        | 56.2        | 65.7        | 32.7        | 1.39        |  |
| XGrad-CAM      | 72.9        | 65.8        | 49.8        | <b>56.2</b> | 69.8        | 33.3        | 1.30        |  | 70.8        | 62.0        | 41.9        | 53.5        | 64.4        | 31.6        | 1.41        |  |
| Layer-CAM      | 73.1        | 66.0        | 50.1        | 55.5        | <b>70.0</b> | 33.0        | 1.29        |  | 70.5        | 61.5        | 28.0        | 54.7        | 65.0        | 32.4        | 1.45        |  |
| ExPerturbation | 73.6        | 66.6        | 37.5        | 44.2        | 64.8        | <b>38.2</b> | 1.59        |  | 74.1        | 66.4        | 37.8        | 43.3        | 62.7        | <b>36.1</b> | 1.74        |  |
| Opti-CAM       | <b>72.2</b> | <b>64.8</b> | 47.3        | 49.2        | 59.4        | 30.5        | 1.34        |  | <b>69.1</b> | <b>59.9</b> | 44.1        | 51.2        | 61.4        | 30.7        | <b>1.34</b> |  |

# Conclusions



# Conclusions

- Classifier specific saliency maps. The probability of prediction of a class is the highest.
- No regularization needed. Optimization only over feature dimensions.

# Conclusions

- Classifier specific saliency maps. The probability of prediction of a class is the highest.
- No regularization needed. Optimization only over feature dimensions.
- Localization is difficult for Opti-CAM. It is mostly human based.

# Conclusions


- Classifier specific saliency maps. The probability of prediction of a class is the highest.
- No regularization needed. Optimization only over feature dimensions.
- Localization is difficult for Opti-CAM. It is mostly human based.
- Average Gain complements the evaluation procedure.



# Preliminaries

## Premise

*The [CLS] token in transformers updates an abstract representation via attention with patch embeddings across the depth of a model. Is it possible to train a similar class representation on CNNs?*



# Preliminaries

## Contributions

- **CAM is Attention**

CAM behaves in a similar manner as attention.

- **Cross Attention: [CLS] on CNNs**

Training of a [CLS] token, alongside standard CNNs.

- **CAM on more architectures**

CAM typical validation only found on ResNet and VGG.

Extension to ConvNeXt.

# CAM is Attention

## Derivation

Let's freshen up:

Asuming  $\mathbf{F}_\ell \in \mathbb{R}^{p_\ell \times d_\ell}$  a reshaping of  $\mathbf{F}_\ell$ , and  $\mathbf{q}_\ell \in \mathbb{R}^{d_\ell}$  the CLS token<sup>3</sup>.

---

<sup>3</sup>  $p$  relates to *width*  $\times$  *height*. Spatial dimension

# CAM is Attention

## Derivation

Let's freshen up:

Asuming  $\mathbf{F}_\ell \in \mathbb{R}^{p_\ell \times d_\ell}$  a reshaping of  $\mathbf{F}_\ell$ , and  $\mathbf{q}_\ell \in \mathbb{R}^{d_\ell}$  the CLS token<sup>3</sup>.

Attention is defined as:

$$\mathbf{a} = \mathbf{A}^\top = \text{softmax} \left( \frac{\mathbf{F}_\ell \mathbf{q}_\ell}{\sqrt{d_\ell}} \right). \quad (8)$$

---

<sup>3</sup>  $p$  relates to *width*  $\times$  *height*. Spatial dimension

# CAM is Attention

## Derivation

Let's freshen up:

Asuming  $\mathbf{F}_\ell \in \mathbb{R}^{p_\ell \times d_\ell}$  a reshaping of  $\mathbf{F}_\ell$ , and  $\mathbf{q}_\ell \in \mathbb{R}^{d_\ell}$  the CLS token<sup>3</sup>.  
Attention is defined as:

$$\mathbf{a} = \mathbf{A}^\top = \text{softmax} \left( \frac{\mathbf{F}_\ell \mathbf{q}_\ell}{\sqrt{d_\ell}} \right). \quad (8)$$

Upon replacing  $\mathbf{q}_\ell$  with an arbitrary vector  $\alpha \in \mathbb{R}^{d_\ell}$

$$\mathbf{a} = h_\ell(\mathbf{F}_\ell \alpha) = h_\ell \left( \sum_k \alpha_k \mathbf{f}_\ell^k \right). \quad (9)$$

With the activation function  $h_\ell(\mathbf{x}) = \text{softmax}(\mathbf{x}/\sqrt{d_\ell})$

---

<sup>3</sup>  $p$  relates to *width*  $\times$  *height*. Spatial dimension

# CAM is Attention

## Pooling or Masking

We intend to integrate attention into a model, such that prediction and explanation are linked.

# CAM is Attention

## Pooling or Masking

We intend to integrate attention into a model, such that prediction and explanation are linked.

**Cross Attention:**

$$\text{CA}_\ell(\mathbf{q}_\ell, F_\ell) := F^\top \mathbf{a} = F_\ell^\top h_\ell(F_\ell \mathbf{q}_\ell) \in \mathbb{R}_\ell^d \quad (10)$$

# CAM is Attention

## Pooling or Masking

We intend to integrate attention into a model, such that prediction and explanation are linked.

**Cross Attention:**

$$\text{CA}_\ell(\mathbf{q}_\ell, F_\ell) := F_\ell^\top \mathbf{a} = F_\ell^\top h_\ell(F_\ell \mathbf{q}_\ell) \in \mathbb{R}_\ell^d \quad (10)$$

Which, reshaping the feature matrix as  $F_\ell^\top = (\phi_\ell^1 \dots \phi_\ell^{p_\ell})$ , a pooled representation can be expressed, via attention vector

$\mathbf{a} = (a_1, \dots, a_{p_\ell})$ :

$$\text{CA}_\ell(\mathbf{q}_\ell, F_\ell) := F_\ell^\top \mathbf{a} = \sum_i a_i \phi_\ell^i. \quad (11)$$

# CAM is Attention

## Pooling or Masking

We intend to integrate attention into a model, such that prediction and explanation are linked.

**Cross Attention:**

$$\text{CA}_\ell(\mathbf{q}_\ell, F_\ell) := F^\top \mathbf{a} = F_\ell^\top h_\ell(F_\ell \mathbf{q}_\ell) \in \mathbb{R}_\ell^d \quad (10)$$

Which, reshaping the feature matrix as  $F_\ell^\top = (\phi_\ell^1 \dots \phi_\ell^{p_\ell})$ , a pooled representation can be expressed, via attention vector

$\mathbf{a} = (a_1, \dots, a_{p_\ell})$ :

$$\text{CA}_\ell(\mathbf{q}_\ell, F_\ell) := F_\ell^\top \mathbf{a} = \sum_i a_i \phi_\ell^i. \quad (11)$$

A pooling similar to GAP, masked in a fashion similar to CAM.

# Cross Attention Stream

Given a model  $f$  we consider locations where critical operations take place:

$$f = g \circ \text{GAP} \circ f_L \circ \dots f_0 \quad (12)$$

# Cross Attention Stream

Given a model  $f$  we consider locations where critical operations take place:

$$f = g \circ \text{GAP} \circ f_L \circ \dots \circ f_0 \quad (12)$$

On another hand, a [CLS] token embedding is initialized, as  $\mathbf{q}_0 \in \mathbb{R}^{d_0}$ , updated with a sequence of embeddings  $\mathbf{q}_\ell \in \mathbb{R}^{d_\ell}$  via interaction along a stream with features on stage  $\ell$ :

$$\mathbf{q}_{\ell+1} = W_\ell \cdot \text{CA}_\ell(\mathbf{q}_\ell, F_\ell) \quad (13)$$

# Cross Attention Stream

Given a model  $f$  we consider locations where critical operations take place:


$$f = g \circ \text{GAP} \circ f_L \circ \dots \circ f_0 \quad (12)$$

On another hand, a [CLS] token embedding is initialized, as  $\mathbf{q}_0 \in \mathbb{R}^{d_0}$ , updated with a sequence of embeddings  $\mathbf{q}_\ell \in \mathbb{R}^{d_\ell}$  via interaction along a stream with features on stage  $\ell$ :

$$\mathbf{q}_{\ell+1} = W_\ell \cdot \text{CA}_\ell(\mathbf{q}_\ell, F_\ell) \quad (13)$$

Visualized as:

Input image  $\mathbf{x}$



# Experimental Set-Up

## Training and models

- Training and validation in train and val sets of ILSVRC 2012 ([15, 16]) respectively.
- Using ResNet 50 original procedure [21]<sup>4</sup>.
- Training ResNet (18-50), ConvNeXt (Small, Base).
- Pascal-VOC 2012[22] and CUB[23] extensions.

---

<sup>4</sup> <https://github.com/pytorch/vision/tree/main/references/classification>

# Experimental Set-Up

## Training and models

- Training and validation in train and val sets of ILSVRC 2012 ([15, 16]) respectively.
- Using ResNet 50 original procedure [21]<sup>4</sup>.
- Training ResNet (18-50), ConvNeXt (Small, Base).
- Pascal-VOC 2012[22] and CUB[23] extensions.


## Evaluation

- Considering the whole image.
- No localization evaluation.
- Raw Attention used for OOD data.
- Ablation on parameters and design choices.

<sup>4</sup> <https://github.com/pytorch/vision/tree/main/references/classification>


# Results

## Qualitative Experiments



# Results

## Qualitative Experiments



# Results

## Quantitative Experiments

### Recognition results

| ACCURACY AND PARAMETERS |      |        |        |        |       |
|-------------------------|------|--------|--------|--------|-------|
| NETWORK                 | POOL | GFLOPS | #PARAM | PARAM% | ACC↑  |
| RESNET-18               | GAP  | 3.648  | 11.69M | 3.71   | 67.28 |
|                         | CA   | 3.652  | 12.13M |        | 67.54 |
| RESNET-50               | GAP  | 8.268  | 25.56M | 27.27  | 74.55 |
|                         | CA   | 8.288  | 32.53M |        | 74.70 |
| CONVNEXT-S              | GAP  | 17.395 | 50.22M | 1.95   | 83.26 |
|                         | CA   | 17.400 | 51.20M |        | 83.14 |
| CONVNEXT-B              | GAP  | 30.747 | 88.59M | 1.96   | 83.72 |
|                         | CA   | 30.753 | 90.33M |        | 83.51 |

# Results

## Quantitative Experiments

### Recognition: Small models

| NETWORK    | METHOD     | POOL      | AD↓                   | AG↑                   | AI↑                   | I↑                    | D↓                    |
|------------|------------|-----------|-----------------------|-----------------------|-----------------------|-----------------------|-----------------------|
| RESNET-18  | Grad-CAM   | GAP       | 17.64                 | 12.73                 | 41.21                 | 63.13                 | <b>10.66</b>          |
|            |            | CA        | <b>16.99</b>          | <b>17.22</b>          | <b>44.95</b>          | <b>65.94</b>          | 10.68                 |
|            | Grad-CAM++ | GAP<br>CA | 19.05<br><b>19.02</b> | 11.16<br><b>14.76</b> | 37.99<br><b>40.82</b> | 62.80<br><b>65.53</b> | <b>10.75</b><br>10.82 |
| CONVNEXT-S | Score-CAM  | GAP       | 13.64                 | 12.98                 | 44.53                 | 62.56                 | <b>11.37</b>          |
|            |            | CA        | <b>11.53</b>          | <b>18.12</b>          | <b>50.32</b>          | <b>65.33</b>          | 11.51                 |
|            | Grad-CAM   | GAP<br>CA | 42.99<br><b>22.09</b> | 1.69<br><b>14.91</b>  | 12.60<br><b>32.65</b> | 48.42<br><b>84.82</b> | <b>30.12</b><br>43.02 |
|            | Grad-CAM++ | GAP       | 56.42                 | 1.32                  | 10.35                 | 48.28                 | <b>33.41</b>          |
|            |            | CA        | <b>51.87</b>          | <b>9.40</b>           | <b>20.55</b>          | <b>84.28</b>          | 52.58                 |
|            | Score-CAM  | GAP<br>CA | 74.79<br><b>64.21</b> | 1.29<br><b>8.81</b>   | 10.10<br><b>18.96</b> | 47.40<br><b>82.92</b> | <b>38.21</b><br>57.46 |

# Results

## Quantitative Experiments

### Recognition: Standard models

| NETWORK    | METHOD     | POOL | AD↓          | AG↑          | AI↑          | I↑           | D↓           |
|------------|------------|------|--------------|--------------|--------------|--------------|--------------|
| RESNET-50  | Grad-CAM   | GAP  | 13.04        | 17.56        | 44.47        | 72.57        | <b>13.24</b> |
|            |            | CA   | <b>12.54</b> | <b>22.67</b> | <b>48.56</b> | <b>75.53</b> | 13.50        |
|            | Grad-CAM++ | GAP  | <b>13.79</b> | 15.87        | 42.08        | 72.32        | <b>13.33</b> |
|            |            | CA   | 13.99        | <b>19.29</b> | <b>44.60</b> | <b>75.21</b> | 13.78        |
|            | Score-CAM  | GAP  | 8.83         | 17.97        | 48.46        | 71.99        | <b>14.31</b> |
|            |            | CA   | <b>7.09</b>  | <b>23.65</b> | <b>54.20</b> | <b>74.91</b> | 14.68        |
| CONVNEXT-B | Grad-CAM   | GAP  | 33.72        | 2.43         | 15.25        | 52.85        | <b>29.57</b> |
|            |            | CA   | <b>19.45</b> | <b>13.96</b> | <b>32.89</b> | <b>86.38</b> | 45.29        |
|            | Grad-CAM++ | GAP  | <b>34.01</b> | 2.37         | 15.60        | 52.83        | <b>29.17</b> |
|            |            | CA   | 36.69        | <b>8.00</b>  | <b>21.95</b> | <b>85.39</b> | 53.42        |
|            | Score-CAM  | GAP  | 43.55        | 2.23         | 15.67        | 50.96        | <b>39.49</b> |
|            |            | CA   | <b>23.51</b> | <b>11.04</b> | <b>27.35</b> | <b>83.41</b> | 60.53        |

## Conclusions

# Conclusions

- CAM behaves like Attention. Attention on final layers ought to behave like CAM.

# Conclusions

- CAM behaves like Attention. Attention on final layers ought to behave like CAM.
- Pooling matters. Smarter pooling yields improved recognition.

# Conclusions

- CAM behaves like Attention. Attention on final layers ought to behave like CAM.
- Pooling matters. Smarter pooling yields improved recognition.
- Deletion requires a revisit. Zeros on images still provide a signal.



# Table of Contents

- Introduction
- 1 Background
- 2 Opti-CAM: Optimizing saliency maps for interpretability
- 3 CA-Stream: Attention-based pooling for interpretable image recognition
- 4 A learning paradigm for interpretable gradients**
- References

# Preliminaries

## Premise

*Gradient information is noisy and hard to interpret, however smoothing it makes it discriminative. Can gradient information be denoised during training, improving interpretability properties?*



# Preliminaries

## Contributions

- Novel regularization for interpretable guided training.
- Improvement of post-hoc interpretations following transparency ideas.

# Interpretable Gradients

## Leveraging Backpropagation

The response ( $y$ ) of a model ( $f$ ) parameters ( $\phi$ ), after forwarding an input  $x$  can be visualized on the input space, via backpropagation:

$$\frac{\partial x}{\partial y} = \frac{\partial x}{\partial F_0} \frac{\partial F_0}{\partial F_1} \dots \frac{\partial F_{-1}}{\partial y} \quad (14)$$



# Interpretable Gradients

## Leveraging Backpropagation

The response ( $y$ ) of a model ( $f$ ) parameters ( $\phi$ ), after forwarding an input  $x$  can be visualized on the input space, via backpropagation:

$$\frac{\partial x}{\partial y} = \frac{\partial x}{\partial F_0} \frac{\partial F_0}{\partial F_1} \dots \frac{\partial F_{-1}}{\partial y} \quad (14)$$

However, due to non-linearities it is noisy:





## Soft Gradients

## Guided Backpropagation [24]

## Smoothgrad [25]

# Interpretable Gradients

## Gradient as Regularization

Goal: Denoise the gradient during training using guided backprop.

*How?* Regularize using guided backprop.

# Interpretable Gradients

## Gradient as Regularization

Goal: Denoise the gradient during training using guided backprop.

*How?* Regularize using guided backprop.

Assuming target labels  $T$ :

$$L_R(X, \theta, T) = \frac{1}{n} \sum_{i=1}^n E(\delta x_i, \delta_G x_i), \quad (15)$$

# Interpretable Gradients

## Gradient as Regularization

Goal: Denoise the gradient during training using guided backprop.

*How?* Regularize using guided backprop.

Assuming target labels  $T$ :


$$L_R(X, \theta, T) = \frac{1}{n} \sum_{i=1}^n E(\delta x_i, \delta_G x_i), \quad (15)$$

And the total loss:

$$L(x, \theta, t) = L_C(x, \theta, t) + \lambda L_R(x, \theta, t), \quad (16)$$

# Interpretable Gradients

## Formalized



# Experimental Set-Up

## Training and models

- Training and validation in train and val sets of CIFAR-100 ([26]) respectively.
- Using an optimized recipe for the dataset <sup>5</sup>.
- Training ResNet (18), MobileNet [27].

---

<sup>5</sup>

# Experimental Set-Up

## Training and models

- Training and validation in train and val sets of CIFAR-100 ([26]) respectively.
- Using an optimized recipe for the dataset <sup>5</sup>.
- Training ResNet (18), MobileNet [27].

## Evaluation

- CIFAR has no center crop. All image it is.
- No localization evaluation.
- Ablation on regularization functions.


---

<sup>5</sup><https://github.com/weiaicunzai/pytorch-cifar100>

# Results

## Qualitative Experiments


### Denoising effect



# Results

## Qualitative Experiments

### CAM Visualizations



# Results

## Quantitative Experiments

### Recognition Metrics

| RECOGNITION METRICS |        |                      |                |
|---------------------|--------|----------------------|----------------|
| MODEL               | ERROR  | $\lambda$            | Acc $\uparrow$ |
| RESNET-18           | -      | -                    | <b>73.42</b>   |
|                     | COSINE | $7.5 \times 10^{-3}$ | 72.86          |
| MOBILENET-V2        | -      | -                    | 59.43          |
|                     | COSINE | $1 \times 10^{-3}$   | <b>62.36</b>   |

# Results

## Quantitative Experiments

### Interpretability metrics

| INTERPRETABLE RECOGNITION METRICS |        |              |              |              |              |              |              |        |              |             |              |              |              |
|-----------------------------------|--------|--------------|--------------|--------------|--------------|--------------|--------------|--------|--------------|-------------|--------------|--------------|--------------|
| RESNET-18                         |        |              |              |              |              |              | MOBILENET-V2 |        |              |             |              |              |              |
| METHOD                            | ERROR  | AD↓          | AG↑          | AI↑          | INS↑         | DEL↓         | METHOD       | ERROR  | AD↓          | AG↑         | AI↑          | INS↑         | DEL↓         |
| GRAD-CAM                          | -      | 30.16        | 15.23        | 29.99        | 58.47        | <b>17.47</b> | GRAD-CAM     | -      | 44.64        | 6.57        | 25.62        | 44.64        | <b>14.34</b> |
|                                   | COSINE | <b>28.09</b> | <b>16.19</b> | <b>31.53</b> | <b>58.76</b> | 17.57        |              | COSINE | <b>40.89</b> | <b>7.31</b> | <b>27.08</b> | <b>45.57</b> | 15.20        |
| GRAD-CAM++                        | -      | 31.40        | 14.17        | 28.47        | 58.61        | <b>17.05</b> | GRAD-CAM++   | -      | 45.98        | 6.12        | 24.10        | 44.72        | <b>14.76</b> |
|                                   | COSINE | <b>29.78</b> | <b>15.07</b> | <b>29.60</b> | <b>58.90</b> | 17.22        |              | COSINE | <b>40.76</b> | <b>6.85</b> | <b>26.46</b> | <b>45.51</b> | 14.92        |
| SCORE-CAM                         | -      | 26.49        | 18.62        | 33.84        | 58.42        | <b>18.31</b> | SCORE-CAM    | -      | 40.55        | 7.85        | 28.57        | 45.62        | <b>14.52</b> |
|                                   | COSINE | <b>24.82</b> | <b>19.49</b> | <b>35.51</b> | <b>59.11</b> | 18.34        |              | COSINE | <b>36.34</b> | <b>9.09</b> | <b>30.50</b> | <b>46.35</b> | 14.72        |
| ABLATION-CAM                      | -      | 31.96        | 14.02        | 28.33        | 58.36        | <b>17.14</b> | ABLATION-CAM | -      | 45.15        | 6.38        | 25.32        | 44.62        | <b>15.03</b> |
|                                   | COSINE | <b>29.90</b> | <b>15.03</b> | <b>29.61</b> | <b>58.70</b> | 17.37        |              | COSINE | <b>41.13</b> | <b>7.03</b> | <b>26.10</b> | <b>45.38</b> | 15.12        |
| AXIOM-CAM                         | -      | 30.16        | 15.23        | 29.98        | 58.47        | <b>17.47</b> | AXIOM-CAM    | -      | 44.65        | 6.57        | 25.62        | 44.64        | 15.27        |
|                                   | COSINE | <b>28.09</b> | <b>16.20</b> | <b>31.53</b> | <b>58.76</b> | 17.57        |              | COSINE | <b>40.89</b> | <b>7.31</b> | <b>27.08</b> | <b>45.57</b> | <b>15.20</b> |

## Conclusions

# Conclusions

- It is possible to denoise CNN gradient while enhancing recognition properties.

# Conclusions

- It is possible to denoise CNN gradient while enhancing recognition properties.
- Negative gradients contribute to stability during training.

# Conclusions

- It is possible to denoise CNN gradient while enhancing recognition properties.
- Negative gradients contribute to stability during training.
- Resolution matters: small inputs require shallow layer targeting.



# References I



J. Malik, P. Arbeláez, J. Carreira, K. Fragkiadaki, R. Girshick, G. Gkioxari, S. Gupta, B. Hariharan, A. Kar, and S. Tulsiani, “The three r’s of computer vision: Recognition, reconstruction and reorganization,” *Pattern Recognition Letters*, vol. 72, pp. 4–14, 2016.



J. McCarthy *et al.*, “What is artificial intelligence,” 2007.



T. Madiega, “Artificial intelligence act,” *European Parliament: European Parliamentary Research Service*, 2021.



Z. C. Lipton, “The mythos of model interpretability: In machine learning, the concept of interpretability is both important and slippery,” *Queue*, vol. 16, no. 3, pp. 31–57, 2018.



K. He, X. Zhang, S. Ren, and J. Sun, “Deep residual learning for image recognition,” in *Proceedings of the IEEE conference on computer vision and pattern recognition*, 2016, pp. 770–778.



J. Devlin, M.-W. Chang, K. Lee, and K. Toutanova, “Bert: Pre-training of deep bidirectional transformers for language understanding,” *arXiv preprint arXiv:1810.04805*, 2018.



A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly *et al.*, “An image is worth 16x16 words: Transformers for image recognition at scale,” *arXiv preprint arXiv:2010.11929*, 2020.



H. Lee, R. Grosse, R. Ranganath, and A. Y. Ng, “Convolutional deep belief networks for scalable unsupervised learning of hierarchical representations,” in *Proceedings of the 26th annual international conference on machine learning*, 2009, pp. 609–616.



A. Chattopadhyay, A. Sarkar, P. Howlader, and V. N. Balasubramanian, “Grad-CAM++: Generalized gradient-based visual explanations for deep convolutional networks,” in *WACV*, 2018.

# References II



P. Vitali, D. Abir, and S. Kate, "Rise: Randomized input sampling for explanation of black-box models," *BMVC*, 2018.



H. Wang, M. Du, F. Yang, and Z. Zhang, "Score-cam: Improved visual explanations via score-weighted class activation mapping," *CoRR*, vol. abs/1910.01279, 2019. [Online]. Available: <http://arxiv.org/abs/1910.01279>



R. C. Fong and A. Vedaldi, "Interpretable explanations of black boxes by meaningful perturbation," in *ICCV*, 2017.



R. Fong, M. Patrick, and A. Vedaldi, "Understanding deep networks via extremal perturbations and smooth masks," in *Proceedings of the IEEE/CVF international conference on computer vision*, 2019, pp. 2950–2958.



D. P. Kingma and J. Ba, "Adam: A method for stochastic optimization," *ICLR*, 2015.



A. Krizhevsky, I. Sutskever, and G. E. Hinton, "Imagenet classification with deep convolutional neural networks," *Advances in neural information processing systems*, vol. 25, 2012.



O. Russakovsky, J. Deng, H. Su, J. Krause, S. Satheesh, S. Ma, Z. Huang, A. Karpathy, A. Khosla, M. Bernstein, A. C. Berg, and L. Fei-Fei, "ImageNet Large Scale Visual Recognition Challenge," *International Journal of Computer Vision (IJCV)*, vol. 115, no. 3, pp. 211–252, 2015.



D. Kermayn, K. Zhang, M. Goldbaum *et al.*, "Labeled optical coherence tomography (oct) and chest x-ray images for classification," *Mendeley data*, vol. 2, no. 2, 2018.



K. Pogorelov, K. R. Randel, C. Griwodz, S. L. Eskeland, T. de Lange, D. Johansen, C. Spampinato, D.-T. Dang-Nguyen, M. Lux, P. T. Schmidt *et al.*, "Kvasir: A multi-class image dataset for computer aided gastrointestinal disease detection," in *Multimedia Systems Conf.*, 2017.



K. Simonyan and A. Zisserman, "Very deep convolutional networks for large-scale image recognition," 2015.

# References III



H. Touvron, M. Cord, A. El-Nouby, P. Bojanowski, A. Joulin, G. Synnaeve, and H. Jégou, "Augmenting convolutional networks with attention-based aggregation," *arXiv preprint arXiv:2112.13692*, 2021.



L. Bottou, "Large-scale machine learning with stochastic gradient descent," in *Proceedings of COMPSTAT'2010: 19th International Conference on Computational Statistics Paris France, August 22-27, 2010 Keynote, Invited and Contributed Papers*. Springer, 2010, pp. 177–186.



M. Everingham, S. M. A. Eslami, L. Van Gool, C. K. I. Williams, J. Winn, and A. Zisserman, "The pascal visual object classes challenge: A retrospective," *International Journal of Computer Vision*, vol. 111, no. 1, pp. 98–136, Jan. 2015.



C. Wah, S. Branson, P. Welinder, P. Perona, and S. Belongie, "Caltech-ucsd birds-200-2011," California Institute of Technology, Tech. Rep. CNS-TR-2011-001, 2011.



J. T. Springenberg, A. Dosovitskiy, T. Brox, and M. Riedmiller, "Striving for simplicity: The all convolutional net," *arXiv preprint arXiv:1412.6806*, 2014.



D. Smilkov, N. Thorat, B. Kim, F. Viégas, and M. Wattenberg, "Smoothgrad: removing noise by adding noise," *arXiv preprint arXiv:1706.03825*, 2017.



A. Krizhevsky, "Learning multiple layers of features from tiny images," pp. 32–33, 2009. [Online]. Available: <https://www.cs.toronto.edu/~kriz/learning-features-2009-TR.pdf>



A. G. Howard, M. Zhu, B. Chen, D. Kalenichenko, W. Wang, T. Weyand, M. Andreetto, and H. Adam, "Mobilenets: Efficient convolutional neural networks for mobile vision applications," *CoRR*, vol. abs/1704.04861, 2017. [Online]. Available: <http://arxiv.org/abs/1704.04861>