

Learning Discriminative Representations to Interpret Image Recognition Models

Thèse de Doctorat

Felipe Torres Figueroa

École Centrale de Marseille

QARMA Laboratoire d'Informatique et de Systèmes (LIS)

Marseille, September 23rd 2024

President: Frédéric Jurie

Reviewers: Frédéric Jurie
Giorgos Tolas

Examiners: Frédéric Precioso
Diane Larlus

Supervisors: Stéphane Ayache
Ronan Sicre
Yannis Avrithis

Table of Contents

- Introduction
- 1 Opti-CAM: Optimizing saliency maps for interpretability
- 2 CA-Stream: Attention-based pooling for interpretable image recognition
- 3 A learning paradigm for interpretable gradients
- Conclusion & Perspectives


Table of Contents

- Introduction
- 1 Opti-CAM: Optimizing saliency maps for interpretability
- 2 CA-Stream: Attention-based pooling for interpretable image recognition
- 3 A learning paradigm for interpretable gradients
- Conclusion & Perspectives

Motivation

Introductory Examples

Consider these situations:



Motivation

Introductory Examples


Consider these situations:




Motivation

Introductory Examples

Consider these situations:



AI accepted in society?



- How do we **know how** a system works?
- How do we **know how** safe a system is?
- If a system fails, **who** is accountable?

We must **explain** the behaviour of these models and their predictions.

Image Recognition Models

An ongoing revolution since 2012





Image Recognition Models

An ongoing revolution since 2012




Contributing factors:

- Widespread of Internet.
- Big Dataset collection.
- Popularization of the GPU.

Image Recognition Models

An ongoing revolution since 2012



Contributing factors:

- Widespread of Internet.
- Big Dataset collection.
- Popularization of the GPU.


Two paradigms:

- Convolutional Neural Networks (CNNs).
- Transformers.

Interpretability

Motivation

We are interested in understanding models, behaving like a black box:



We want to *know why* $f(x) \rightarrow y$

Interpretability

A paradigm on interpretable studies

Interpretability approaches have been categorized by two studies [1, 2]:

1. Passive or Active?


*To add or not modifications
to the model.*



A paradigm on interpretable studies

1. **Passive or Active?**
To add or not modifications to the model.

- ## 2. Type of Explanation
- Is it an example? An Attribution? Hidden Semantics? Or a Rule?*



Interpretability

A paradigm on interpretable studies

Interpretability approaches have been categorized by two studies [1, 2]:

1. **Passive or Active?**¹


To add or not modifications to the model.

2. **Type of Explanation**

Is it an example? An Attribution? Hidden Semantics? Or a Rule?

3. **Local or Global?**

Is it a brief explanation? Does it cover more data/ of the model? Does it explain the model/data entirely?




¹ Lipton defines Interpretability Similarly.
Transparency or Post-hoc Interpretations

Interpretability

Post-hoc Interpretability


Post-hoc Interpretability



Interpretability


Post-hoc Interpretability

Post-hoc Interpretability




Interpretability

Post-hoc Interpretability




Interpretability

Post-hoc Interpretability




Interpretability

Post-hoc Interpretability



Interpretability

Post-hoc Interpretability



Contributions

High Level

We are interested in understanding models, behaving like a black box:



How to *know why* $f(x) \rightarrow y$

Contributions

High Level

We are interested in understanding models,
behaving like a black box:



How to *know why* $f(x) \rightarrow y$

Saliency Maps



Contributions

High Level

We are interested in understanding models, behaving like a black box:



How to *know why* $f(x) \rightarrow y$

Saliency Maps



Gradients



Contributions

High Level

We are interested in understanding models, behaving like a black box:



How to *know why* $f(x) \rightarrow y$

Saliency Maps



Gradients



Other Representations



Contributions

High Level

We are interested in understanding models, behaving like a black box:



How to *know why* $f(x) \rightarrow y$

Saliency Maps



Gradients



Other Representations



We demonstrate improvements in these approaches

Contributions

Table of Contents

- Introduction
- 1 Opti-CAM: Optimizing saliency maps for interpretability
- 2 CA-Stream: Attention-based pooling for interpretable image recognition
- 3 A learning paradigm for interpretable gradients
- Conclusion & Perspectives

Preliminaries

Premise

Is it possible to optimize a saliency map probability with the highest probability of prediction?



$$p(x)_{774} = 0.745$$



$$p(x)_{774}^{CAM} = 0.728$$

Premise


Is it possible to optimize a saliency map probability with the highest probability of prediction?



$$p(x)_{774} = 0.745$$




$$p(x)_{774}^{CAM} = 0.728$$



$$p(x)_{774}^{Optimal} = 0.853$$


Background

CAM Methods




Background

CAM Methods



Background

CAM Methods



Background

CAM Formalized

We define CAM according to the following equation:

$$S_{\ell}^c(\mathbf{u}) := h \left(\sum_k w_k^c A_l^k \right) \quad (1)$$

Background

CAM Formalized

We define CAM according to the following equation:

$$S_{\ell}^c(\mathbf{u}) := h \left(\sum_k w_k^c A_l^k \right) \quad (1)$$

The weighting coefficient $\mathbf{w}_k^c$ defining the CAM variant.

Background

CAM Formalized

We define CAM according to the following equation:

$$S_{\ell}^c(\mathbf{u}) := h \left(\sum_k w_k^c A_l^k \right) \quad (1)$$

The weighting coefficient $\mathbf{w}_k^c$ defining the CAM variant.
Milestone CAM models:

CAM[3]:

w_k^c comes from the weights
for one class. Only on the
last convolutional layers
followed by GAP.

Background

CAM Formalized

We define CAM according to the following equation:

$$S_{\ell}^c(\mathbf{u}) := h \left(\sum_k w_k^c A_l^k \right) \quad (1)$$

The weighting coefficient w_k^c defining the CAM variant.
Milestone CAM models:

CAM[3]:

w_k^c comes from the weights for one class. Only on the last convolutional layers followed by GAP.

Grad-CAM/Grad-CAM++

[4, 5]:

w_k^c is computed using gradients flowing from prediction to input space. It can be used at any point of the model.

Background

CAM Formalized

We define CAM according to the following equation:

$$S_{\ell}^c(\mathbf{u}) := h \left(\sum_k w_k^c A_l^k \right) \quad (1)$$

The weighting coefficient $\mathbf{w}_k^c$ defining the CAM variant.
Milestone CAM models:

CAM[3]:

w_k^c comes from the weights for one class. Only on the last convolutional layers followed by GAP.

Grad-CAM/Grad-CAM++ [4, 5]:

w_k^c is computed using gradients flowing from prediction to input space. It can be used at any point of the model.

Score-CAM[6]:

w_k^c is computed from the contributions of individual feature maps masking the input image.

Background

CAM Formalized

We define CAM according to the following equation:

$$S_{\ell}^c(\mathbf{u}) := h \left(\sum_k w_k^c A_l^k \right) \quad (1)$$

The weighting coefficient $\mathbf{w}_k^c$ defining the CAM variant.
Milestone CAM models:

CAM[3]:

w_k^c comes from the weights for one class. Only on the last convolutional layers followed by GAP.

Grad-CAM/Grad-CAM++ [4, 5]:

w_k^c is computed using gradients flowing from prediction to input space. It can be used at any point of the model.


Score-CAM[6]:

w_k^c is computed from the contributions of individual feature maps masking the input image.

The approach found in Score-CAM is interesting. Diving deeper:

Background

Score-CAM



Background

Score-CAM Formalized

Following equation (1):

Background

Score-CAM Formalized

Following equation (1):

- $h := \text{ReLU}$
- $\mathbf{w}_k^c := \text{softmax}(\mathbf{u}^c)_k$

Background

Score-CAM Formalized

Following equation (1):


- $h := \text{ReLU}$
- $\mathbf{w}_k^c := \text{softmax}(\mathbf{u}^c)_k$

With feature map contribution $\mathbf{u}^c_k$:

$$u_k^c := f(\mathbf{x} \odot n(\text{up}(A_\ell^k)))_c - f(\mathbf{x})_c \quad (2)$$


Opti-CAM

Approach



Opti-CAM

Approach



Consider S_ℓ , a function of $\mathbf{u}$ and $\mathbf{x}$:

$$S_\ell(\mathbf{x}; \mathbf{u}) := \sum_k \text{softmax}(\mathbf{u})_k A_\ell^k \quad (3)$$

$\mathbf{u}^*$ is the vector that maximizes class c probability, when masking the the input with S_ℓ .

$$\mathbf{u}^* := \underset{\mathbf{u}}{\operatorname{argmax}} F_\ell^c(\mathbf{x}; \mathbf{u}) \quad (4)$$

And the objective $F_\ell^c(\mathbf{x}; \mathbf{u}) := f(\mathbf{x} \odot n(\text{up}(S_\ell(\mathbf{x}; \mathbf{u}))))$

Evaluating Interpretability

Three paradigms

Evaluating Interpretability

Three paradigms

- **Interpretable Image Recognition[5]**

Average Drop (AD), Average Increase (AI), Average Gain (AG).

Evaluating Interpretability

Three paradigms

- **Interpretable Image Recognition[5]**
Average Drop (AD), Average Increase (AI), Average Gain (AG).
- **Causality Analysis[7]**
Insertion & Deletion.
- **Weakly Supervised Localization Approach**

Evaluating Interpretability

Three paradigms

- **Interpretable Image Recognition[5]**

Average Drop (AD), Average Increase (AI), Average Gain (AG).

- **Causality Analysis[7]**


Insertion & Deletion.

- **Weakly Supervised Localization Approach**

Official Metric (OM), Localization Error (LE), Pixel-wise F_1 score (F1) Box Accuracy (BA), Standard Pointing Game (SP), Energy Pointing Game (EP).


Results

Qualitative Evaluation



Results

Qualitative Evaluation



Results

Quantitative Evaluation

Classification Experiments:

METHOD	RESNET50						VGG16					
	AD ↓	AG ↑	AI ↑	I ↑	D ↓	T(s)	AD ↓	AG ↑	AI ↑	I ↑	D ↓	T(s)
Fake-CAM	0.8	1.6	46.0	50.7	28.1	0.00	0.5	0.6	42.6	46.1	26.9	0.00
Grad-CAM	12.2	17.6	44.4	66.3	14.7	0.03	14.2	14.7	40.6	64.1	11.6	0.02
Grad-CAM++	12.9	16.0	42.1	66.0	14.7	0.03	17.1	10.2	33.4	62.9	12.2	0.02
Score-CAM	8.6	26.6	56.7	65.7	16.3	15.22	13.5	15.6	41.7	62.5	12.1	3.11
Ablation-CAM	12.5	16.4	42.8	65.9	14.6	18.26	15.5	12.6	36.9	63.8	11.4	2.98
XGrad-CAM	12.2	17.6	44.4	66.3	14.7	0.03	13.8	14.8	41.2	64.1	11.7	0.02
Layer-CAM	15.6	15.0	38.8	67.0	14.2	0.08	48.9	3.1	13.5	58.3	6.4	0.07
ExPerturbation	38.1	9.5	22.5	70.7	15.0	152.96	43.0	7.1	20.5	61.1	15.0	83.20
Opti-CAM	1.5	68.8	92.8	62.0	19.7	4.15	1.3	71.2	92.7	59.2	11.0	3.94

Results

Quantitative Evaluation

Classification Experiments:

METHOD	RESNET50						VGG16					
	AD ↓	AG ↑	AI ↑	I ↑	D ↓	T(s)	AD ↓	AG ↑	AI ↑	I ↑	D ↓	T(s)
Fake-CAM	0.8	1.6	46.0	50.7	28.1	0.00	0.5	0.6	42.6	46.1	26.9	0.00
Grad-CAM	12.2	17.6	44.4	66.3	14.7	0.03	14.2	14.7	40.6	64.1	11.6	0.02
Grad-CAM++	12.9	16.0	42.1	66.0	14.7	0.03	17.1	10.2	33.4	62.9	12.2	0.02
Score-CAM	8.6	26.6	56.7	65.7	16.3	15.22	13.5	15.6	41.7	62.5	12.1	3.11
Ablation-CAM	12.5	16.4	42.8	65.9	14.6	18.26	15.5	12.6	36.9	63.8	11.4	2.98
XGrad-CAM	12.2	17.6	44.4	66.3	14.7	0.03	13.8	14.8	41.2	64.1	11.7	0.02
Layer-CAM	15.6	15.0	38.8	67.0	14.2	0.08	48.9	3.1	13.5	58.3	6.4	0.07
ExPerturbation	38.1	9.5	22.5	70.7	15.0	152.96	43.0	7.1	20.5	61.1	15.0	83.20
Opti-CAM	1.5	68.8	92.8	62.0	19.7	4.15	1.3	71.2	92.7	59.2	11.0	3.94

The optimization objective is well aligned with AD, AI, AG.
Insertion and Deletion are better suited for sparse saliency maps.

Results

Quantitative Evaluation

Localization Experiments:

METHOD	RESNET50								VGG16							
	OM↓	LE↓	F1↑	BA↑	SP↑	EP↑	SM↓		OM↓	LE↓	F1↑	BA↑	SP↑	EP↑	SM↓	
Fake-CAM	63.6	54.0	57.7	47.9	99.8	28.5	0.98		64.7	54.0	57.7	47.9	99.8	28.5	1.07	
Grad-CAM	72.9	65.8	49.8	56.2	69.8	33.3	1.30		71.1	62.3	42.0	54.2	64.8	32.0	1.39	
Grad-CAM++	73.1	66.1	50.4	56.2	69.9	33.1	1.29		70.8	61.9	44.3	55.2	66.2	32.3	1.38	
Score-CAM	72.2	64.9	49.6	54.5	68.7	32.4	1.25		71.2	62.5	45.3	58.5	68.2	33.4	1.40	
Ablation-CAM	72.8	65.7	50.2	56.1	69.9	33.1	1.26		71.3	62.6	43.2	56.2	65.7	32.7	1.39	
XGrad-CAM	72.9	65.8	49.8	56.2	69.8	33.3	1.30		70.8	62.0	41.9	53.5	64.4	31.6	1.41	
Layer-CAM	73.1	66.0	50.1	55.5	70.0	33.0	1.29		70.5	61.5	28.0	54.7	65.0	32.4	1.45	
ExPerturbation	73.6	66.6	37.5	44.2	64.8	38.2	1.59		74.1	66.4	37.8	43.3	62.7	36.1	1.74	
Opti-CAM	72.2	64.8	47.3	49.2	59.4	30.5	1.34		69.1	59.9	44.1	51.2	61.4	30.7	1.34	

Results

Quantitative Evaluation

Localization Experiments:

METHOD	RESNET50								VGG16							
	OM↓	LE↓	F1↑	BA↑	SP↑	EP↑	SM↓		OM↓	LE↓	F1↑	BA↑	SP↑	EP↑	SM↓	
Fake-CAM	63.6	54.0	57.7	47.9	99.8	28.5	0.98		64.7	54.0	57.7	47.9	99.8	28.5	1.07	
Grad-CAM	72.9	65.8	49.8	56.2	69.8	33.3	1.30		71.1	62.3	42.0	54.2	64.8	32.0	1.39	
Grad-CAM++	73.1	66.1	50.4	56.2	69.9	33.1	1.29		70.8	61.9	44.3	55.2	66.2	32.3	1.38	
Score-CAM	72.2	64.9	49.6	54.5	68.7	32.4	1.25		71.2	62.5	45.3	58.5	68.2	33.4	1.40	
Ablation-CAM	72.8	65.7	50.2	56.1	69.9	33.1	1.26		71.3	62.6	43.2	56.2	65.7	32.7	1.39	
XGrad-CAM	72.9	65.8	49.8	56.2	69.8	33.3	1.30		70.8	62.0	41.9	53.5	64.4	31.6	1.41	
Layer-CAM	73.1	66.0	50.1	55.5	70.0	33.0	1.29		70.5	61.5	28.0	54.7	65.0	32.4	1.45	
ExPerturbation	73.6	66.6	37.5	44.2	64.8	38.2	1.59		74.1	66.4	37.8	43.3	62.7	36.1	1.74	
Opti-CAM	72.2	64.8	47.3	49.2	59.4	30.5	1.34		69.1	59.9	44.1	51.2	61.4	30.7	1.34	

Localization is hard. It acts more aligned to human interpretations than model ones.

Conclusions

- Classifier specific saliency maps. The probability of prediction of a class is the highest.
- No regularization needed unlike [8, 9]. Optimization only over feature dimensions.
- Localization is difficult for Opti-CAM. It is mostly human based.
- Average Gain complements the evaluation procedure.
- Published at Computer Vision and Image Understanding (CVIU) 2024.

Table of Contents

- Introduction
- 1 ● Opti-CAM: Optimizing saliency maps for interpretability
- 2 ● **CA-Stream: Attention-based pooling for interpretable image recognition**
- 3 ● A learning paradigm for interpretable gradients
- Conclusion & Perspectives

Preliminaries

Premise

Is it possible to incorporate transformers ideas onto CNNs, enhancing the predictive power of CAM explanations?



$$p = 0.7623$$



$$p = 0.7145$$

Preliminaries

Premise

Is it possible to incorporate transformers ideas onto CNNs, enhancing the predictive power of CAM explanations?



$$p = 0.7623$$



$$p = 0.7145$$



$$p = 0.7478$$

Background


Self-attention is the building block of Transformers:

Background

Self Attention

Self-attention is the building block of Transformers:

$$\text{Self Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_e}} \right) V \quad (5)$$




Background

Self Attention

Self-attention is the building block of Transformers:

$$\text{Self Attention}(Q, K, V) = \text{softmax} \left(\frac{QK^T}{\sqrt{d_e}} \right) V \quad (5)$$



Computation of the relevance of each element relative to others in an embedding.

CAM is Attention

Derivation

From equation (5), the equation matrix $\mathbf{a}$:

$$\mathbf{a} = \text{softmax} \left(\frac{Q \times K^T}{\sqrt{d_\ell}} \right) \quad (6)$$

CAM is Attention

Derivation

From equation (5), the equation matrix $\mathbf{a}$:

$$\mathbf{a} = \text{softmax} \left(\frac{Q \times K^T}{\sqrt{d_\ell}} \right) \quad (6)$$

Now, assuming an arbitrary vector α to compute attention with, and embedding V a feature map A_ℓ :

CAM is Attention

Derivation

From equation (5), the equation matrix $\mathbf{a}$:

$$\mathbf{a} = \text{softmax} \left(\frac{Q \times K^T}{\sqrt{d_\ell}} \right) \quad (6)$$

Now, assuming an arbitrary vector α to compute attention with, and embedding V a feature map A_ℓ :

$$\mathbf{a} = \text{softmax} \left(\frac{\alpha \times A_\ell}{\sqrt{d_\ell}} \right) \quad (7)$$

CAM is Attention

Derivation

From equation (5), the equation matrix $\mathbf{a}$:

$$\mathbf{a} = \text{softmax} \left(\frac{Q \times K^T}{\sqrt{d_\ell}} \right) \quad (6)$$

Now, assuming an arbitrary vector α to compute attention with, and embedding V a feature map A_ℓ :

$$\mathbf{a} = \text{softmax} \left(\frac{\alpha \times A_\ell}{\sqrt{d_\ell}} \right) \quad (7)$$

Where $\text{softmax}(\dots/\sqrt{d_\ell})$ is h_ℓ .

CAM is Attention

Derivation

From equation (5), the equation matrix $\mathbf{a}$:

$$\mathbf{a} = \text{softmax} \left(\frac{Q \times K^T}{\sqrt{d_\ell}} \right) \quad (6)$$

Now, assuming an arbitrary vector α to compute attention with, and embedding V a feature map A_ℓ :

$$\mathbf{a} = \text{softmax} \left(\frac{\alpha \times A_\ell}{\sqrt{d_\ell}} \right) \quad (7)$$

Where $\text{softmax}(\dots/\sqrt{d_\ell})$ is h_ℓ . Thus

$$\mathbf{a} = h_\ell(\alpha \times A_\ell) = h_\ell \left(\sum_k \alpha_k A_\ell^k \right) \quad (8)$$

CAM is Attention

Pooling or Masking

We intend to integrate attention into a model, such that prediction and explanation are linked.

CAM is Attention

Pooling or Masking

We intend to integrate attention into a model, such that prediction and explanation are linked.

Cross Attention: Assuming an arbitrary embedding $\mathbf{q}$:

$$\text{CA}_\ell(\mathbf{q}_\ell, A_\ell) := A^\top \mathbf{a} = A_\ell^\top h_\ell(A_\ell \mathbf{q}_\ell), \in \mathbb{R}_\ell^d \quad (9)$$

CAM is Attention

Pooling or Masking

We intend to integrate attention into a model, such that prediction and explanation are linked.

Cross Attention: Assuming an arbitrary embedding $\mathbf{q}$:

$$\text{CA}_\ell(\mathbf{q}_\ell, A_\ell) := A^\top \mathbf{a} = A_\ell^\top h_\ell(A_\ell \mathbf{q}_\ell), \in \mathbb{R}_\ell^d \quad (9)$$

With a flattened A_ℓ , a pooled representation can be expressed via attention vector $\mathbf{a} = (a_1, \dots, a_{p_\ell})$:

$$\text{CA}_\ell(\mathbf{q}_\ell, A_\ell) := A_\ell^\top \mathbf{a} \quad (10)$$

CAM is Attention

Pooling or Masking

We intend to integrate attention into a model, such that prediction and explanation are linked.

Cross Attention: Assuming an arbitrary embedding $\mathbf{q}$:

$$\text{CA}_\ell(\mathbf{q}_\ell, A_\ell) := A^\top \mathbf{a} = A_\ell^\top h_\ell(A_\ell \mathbf{q}_\ell), \in \mathbb{R}_\ell^d \quad (9)$$

With a flattened A_ℓ , a pooled representation can be expressed via attention vector $\mathbf{a} = (a_1, \dots, a_{p_\ell})$:

$$\text{CA}_\ell(\mathbf{q}_\ell, A_\ell) := A_\ell^\top \mathbf{a} \quad (10)$$

A pooling similar to GAP, masked in a fashion similar to CAM.

Cross Attention Stream

Given a model f we consider locations where critical operations take place:

$$f = g \circ \text{GAP} \circ f_L \circ \dots f_0 \quad (11)$$

Cross Attention Stream

Given a model f we consider locations where critical operations take place:

$$f = g \circ \text{GAP} \circ f_L \circ \dots \circ f_0 \quad (11)$$

On another hand, a [CLS] token embedding is initialized, as $\mathbf{q}_0 \in \mathbb{R}^{d_0}$, updated with a sequence of embeddings $\mathbf{q}_\ell \in \mathbb{R}^{d_\ell}$ via interaction along a stream with features on stage ℓ :

$$\mathbf{q}_{\ell+1} = W_\ell \cdot \text{CA}_\ell(\mathbf{q}_\ell, A_\ell) \quad (12)$$

Cross Attention Stream

Given a model f we consider locations where critical operations take place:


$$f = g \circ \text{GAP} \circ f_L \circ \dots \circ f_0 \quad (11)$$

On another hand, a [CLS] token embedding is initialized, as $\mathbf{q}_0 \in \mathbb{R}^{d_0}$, updated with a sequence of embeddings $\mathbf{q}_\ell \in \mathbb{R}^{d_\ell}$ via interaction along a stream with features on stage ℓ :

$$\mathbf{q}_{\ell+1} = W_\ell \cdot \text{CA}_\ell(\mathbf{q}_\ell, A_\ell) \quad (12)$$


Visualized as:

Input image $\mathbf{x}$



Results

Qualitative Experiments



Results

Quantitative Experiments

Recognition results

ACCURACY AND PARAMETERS					
NETWORK	POOL	GFLOPS	#PARAM	PARAM%	ACC↑
RESNET-18	GAP	3.648	11.69M	3.71	67.28
	CA	3.652	12.13M		67.54
RESNET-50	GAP	8.268	25.56M	27.27	74.55
	CA	8.288	32.53M		74.70
CONVNEXT-S	GAP	17.395	50.22M	1.95	83.26
	CA	17.400	51.20M		83.14
CONVNEXT-B	GAP	30.747	88.59M	1.96	83.72
	CA	30.753	90.33M		83.51

Results

Quantitative Experiments

NETWORK	METHOD	POOL	AD↓	AG↑	AI↑	I↑	D↓
RESNET-50	Grad-CAM	GAP	13.04	17.56	44.47	72.57	13.24
		CA	12.54	22.67	48.56	75.53	13.50
	Grad-CAM++	GAP	13.79	15.87	42.08	72.32	13.33
		CA	13.99	19.29	44.60	75.21	13.78
	Score-CAM	GAP	8.83	17.97	48.46	71.99	14.31
		CA	7.09	23.65	54.20	74.91	14.68
CONVNEXT-B	Grad-CAM	GAP	33.72	2.43	15.25	52.85	29.57
		CA	19.45	13.96	32.89	86.38	45.29
	Grad-CAM++	GAP	34.01	2.37	15.60	52.83	29.17
		CA	36.69	8.00	21.95	85.39	53.42
	Score-CAM	GAP	43.55	2.23	15.67	50.96	39.49
		CA	23.51	11.04	27.35	83.41	60.53

Conclusions

- CAM behaves like Attention. Attention on final layers ought to behave like CAM.
- Pooling matters. Smarter pooling yields improved interpretations, maintaining recognition performance.
- Deletion requires a revisit. Zeros on images still provide a signal.
- Published in the XAI Workshop at CVPR 2024.

Table of Contents

- Introduction
- 1 Opti-CAM: Optimizing saliency maps for interpretability
- 2 CA-Stream: Attention-based pooling for interpretable image recognition
- 3 A learning paradigm for interpretable gradients**
- Conclusion & Perspectives

Background

Leveraging Backpropagation

The response (y) of a model (f) parameters (ϕ), after forwarding an input x can be visualized on the input space, via backpropagation:

$$\frac{\partial x}{\partial y} = \frac{\partial x}{\partial F_0} \frac{\partial F_0}{\partial F_1} \dots \frac{\partial F_{-1}}{\partial y} \quad (13)$$


Background

Soft Gradients

Soft gradients can be achieved by two approaches:

Guided Backpropagation [10]

Achieved by constraining backprop to only let positive gradients flow:




Background

Soft Gradients

Soft gradients can be achieved by two approaches:


Guided Backpropagation [10]

Achieved by constraining backprop to only let positive gradients flow:



Smoothgrad [11]

Achieved by iteratively adding noise to the input space prior to forward pass:



Interpretable Gradients


Formalized

Goal: Denoise the gradient during training using guided backprop.
How? Regularize using guided backprop.

Interpretable Gradients

Formalized

Goal: Denoise the gradient during training using guided backprop.
How? Regularize using guided backprop.



Interpretable Gradients

Smooth Gradient as Regularization

Assuming target labels T :

$$L_R(X, \theta, T) = \frac{1}{n} \sum_{i=1}^n E(\delta x_i, \delta_G x_i), \quad (14)$$

Interpretable Gradients

Smooth Gradient as Regularization

Assuming target labels T :

$$L_R(X, \theta, T) = \frac{1}{n} \sum_{i=1}^n E(\delta x_i, \delta_G x_i), \quad (14)$$


And the total loss:

$$L(x, \theta, t) = L_C(x, \theta, t) + \lambda L_R(x, \theta, t), \quad (15)$$

Results

Qualitative Experiments


Denoising effect



Results

Qualitative Experiments

CAM Visualizations



Results

Quantitative Experiments

Recognition Metrics

RECOGNITION METRICS			
MODEL	ERROR	λ	ACC $\uparrow$
RESNET-18	-	-	73.42
	COSINE	7.5×10^{-3}	72.86
MOBILENET-V2	-	-	59.43
	COSINE	1×10^{-3}	62.36

Results

Quantitative Experiments

Interpretability metrics

INTERPRETABLE RECOGNITION METRICS						
RESNET-18						
METHOD	λ	AD↓	AG↑	AI↑	INS↑	DEL↓
GRAD-CAM	-	30.16	15.23	29.99	58.47	17.47
	COSINE	28.09	16.19	31.53	58.76	17.57
GRAD-CAM++	-	31.40	14.17	28.47	58.61	17.05
	COSINE	29.78	15.07	29.60	58.90	17.22
SCORE-CAM	-	26.49	18.62	33.84	58.42	18.31
	COSINE	24.82	19.49	35.51	59.11	18.34
ABLATION-CAM	-	31.96	14.02	28.33	58.36	17.14
	COSINE	29.90	15.03	29.61	58.70	17.37
AXIOM-CAM	-	30.16	15.23	29.98	58.47	17.47
	COSINE	28.09	16.20	31.53	58.76	17.57

Conclusions

- Denoising CNN gradients improve interpretability properties, maintaining recognition.
- Negative gradients contribute to stability during training.
- Resolution matters: small inputs require shallow layer targeting.
- Published in VISAPP 2024.

Table of Contents

- Introduction
- 1 Opti-CAM: Optimizing saliency maps for interpretability
- 2 CA-Stream: Attention-based pooling for interpretable image recognition
- 3 A learning paradigm for interpretable gradients
- Conclusion & Perspectives

Table of Contents

- Introduction
- 1 Opti-CAM: Optimizing saliency maps for interpretability
- 2 CA-Stream: Attention-based pooling for interpretable image recognition
- 3 A learning paradigm for interpretable gradients
- Conclusion & Perspectives

References I



Z. C. Lipton, “The mythos of model interpretability: In machine learning, the concept of interpretability is both important and slippery.” *Queue*, vol. 16, no. 3, pp. 31–57, 2018.



Y. Zhang, P. Tiño, A. Leonardis, and K. Tang, “A survey on neural network interpretability,” *IEEE Transactions on Emerging Topics in Computational Intelligence*, vol. 5, no. 5, pp. 726–742, 2021.



B. Zhou, A. Khosla, A. Lapedriza, A. Oliva, and A. Torralba, “Learning deep features for discriminative localization,” in *CVPR*, 2016.



R. R. Selvaraju, A. Das, R. Vedantam, M. Cogswell, D. Parikh, and D. Batra, “Grad-cam: Why did you say that? visual explanations from deep networks via gradient-based localization,” *CoRR*, vol. abs/1610.02391, 2016. [Online]. Available: <http://arxiv.org/abs/1610.02391>



A. Chattopadhyay, A. Sarkar, P. Howlader, and V. N. Balasubramanian, “Grad-CAM++: Generalized gradient-based visual explanations for deep convolutional networks,” in *WACV*, 2018.



H. Wang, M. Du, F. Yang, and Z. Zhang, “Score-cam: Improved visual explanations via score-weighted class activation mapping,” *CoRR*, vol. abs/1910.01279, 2019. [Online]. Available: <http://arxiv.org/abs/1910.01279>



P. Vitali, D. Abir, and S. Kate, “Rise: Randomized input sampling for explanation of black-box models,” *BMVC*, 2018.



R. C. Fong and A. Vedaldi, “Interpretable explanations of black boxes by meaningful perturbation,” in *ICCV*, 2017.



R. Fong, M. Patrick, and A. Vedaldi, “Understanding deep networks via extremal perturbations and smooth masks,” in *Proceedings of the IEEE/CVF international conference on computer vision*, 2019, pp. 2950–2958.

References II



J. T. Springenberg, A. Dosovitskiy, T. Brox, and M. Riedmiller, "Striving for simplicity: The all convolutional net," *arXiv preprint arXiv:1412.6806*, 2014.



D. Smilkov, N. Thorat, B. Kim, F. Viégas, and M. Wattenberg, "Smoothgrad: removing noise by adding noise," *arXiv preprint arXiv:1706.03825*, 2017.